

Red wine quality prediction

Machine Learning course project

Nicolò Rubattu

nicolo.rubattu@student.supsi.ch

The Dataset

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
0	7.4	0.7	0.0	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	5



- Red and White Wine

The inputs include objective tests (e.g. PH values) and the output is based on sensory data (median of at least 3 evaluations made by wine experts). Each expert graded the wine quality between 0 (very bad) and 10 (very excellent)

- Objective: predict quality → *Regression + Classification*


- *Credits:*

P. Cortez, A. Cerdeira, F. Almeida, T. Matos and J. Reis.
Modeling wine preferences by data mining from physicochemical properties.
In Decision Support Systems, Elsevier, 47(4):547-553. ISSN: 0167-9236.

Available at: [Elsevier] <http://dx.doi.org/10.1016/j.dss.2009.05.016>
[Pre-press (pdf)] <http://www3.dsi.uminho.pt/pcortez/winequality09.pdf>
[bib] <http://www3.dsi.uminho.pt/pcortez/dss09.bib>

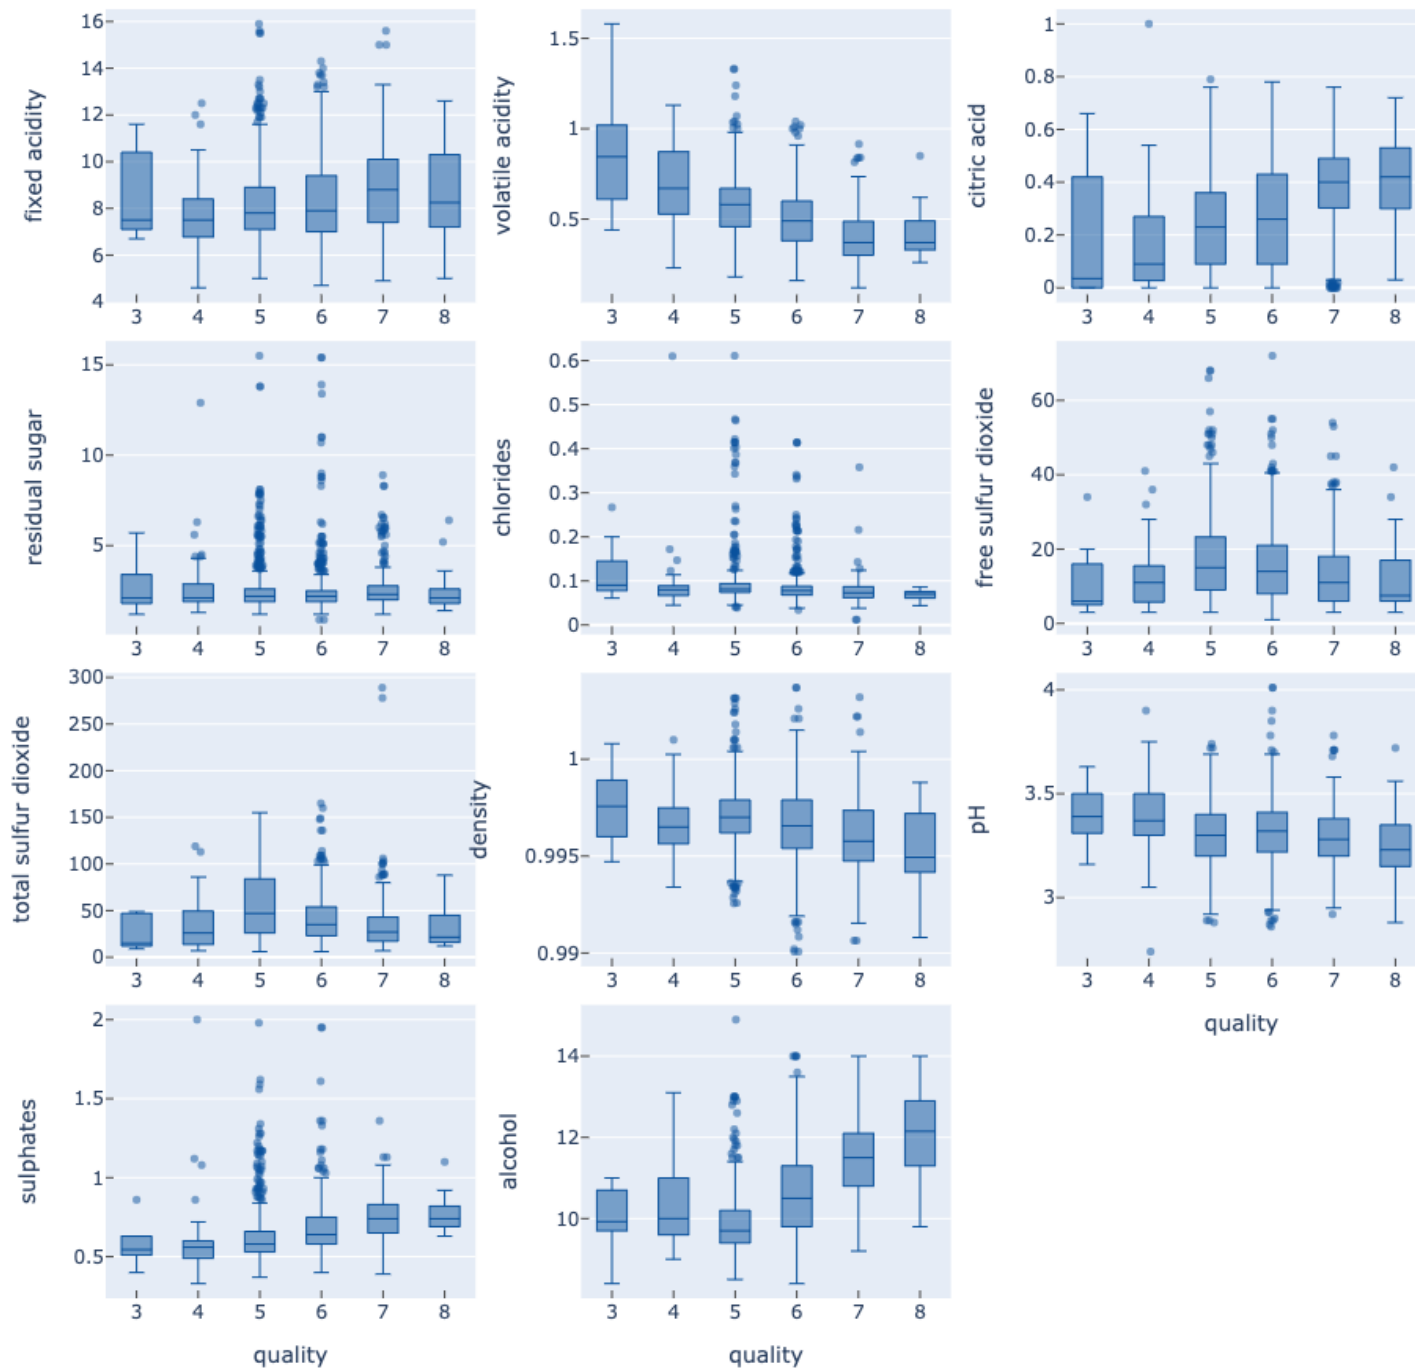
Data Analysis

- Number of Instances: red wine - 1599; white wine 4898
- Number of Attributes: 11 (fixed acidity, volatile acidity, citric acid, residual sugar, chlorides, free sulfur dioxide, total sulfur dioxide, density, pH, sulphates, alcohol) + quality
- Statistic analysis, outliers, NaN values, correlation analysis




Red Wine choice

BoxPlot overview of features and target correlation with outliers check.

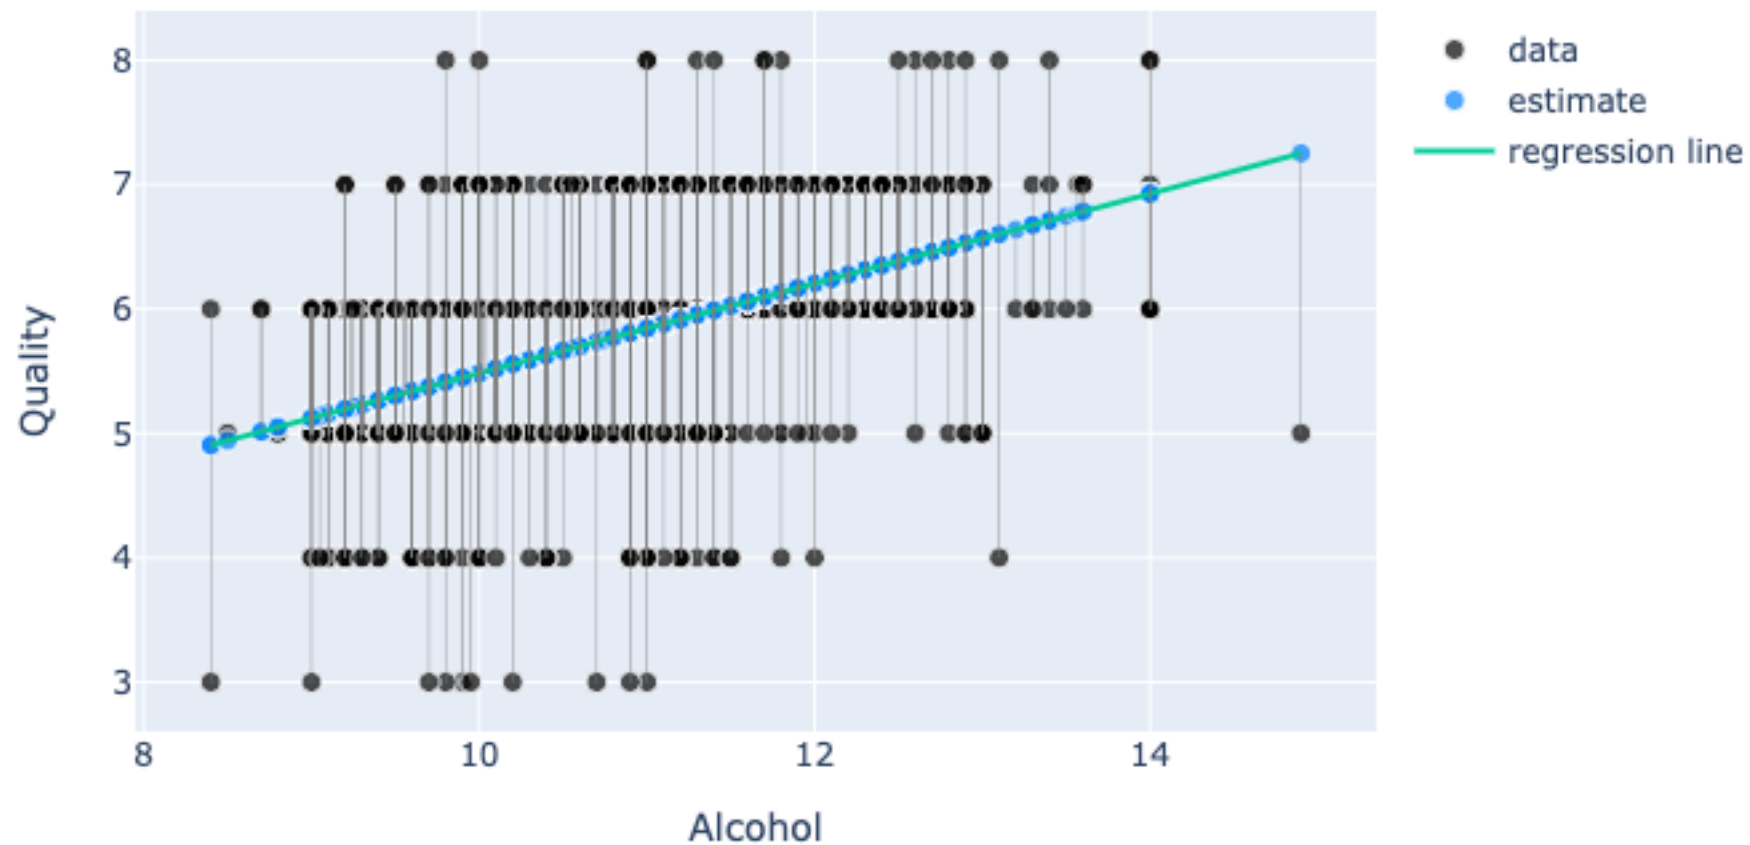


Quality distribution



Regression

Univariate linear regression (alcohol)




MSE
0.504

R^2
0.2267

Regression

Simple Multivariate linear regression (alcohol + volatile acidity)



MSE

0.414

Adjusted R²

0.2723

Regression

Multivariate linear regression (all features) with Stratified KFold Validation

- Training Set

MSE	Adjusted R ²
0.416	0.357

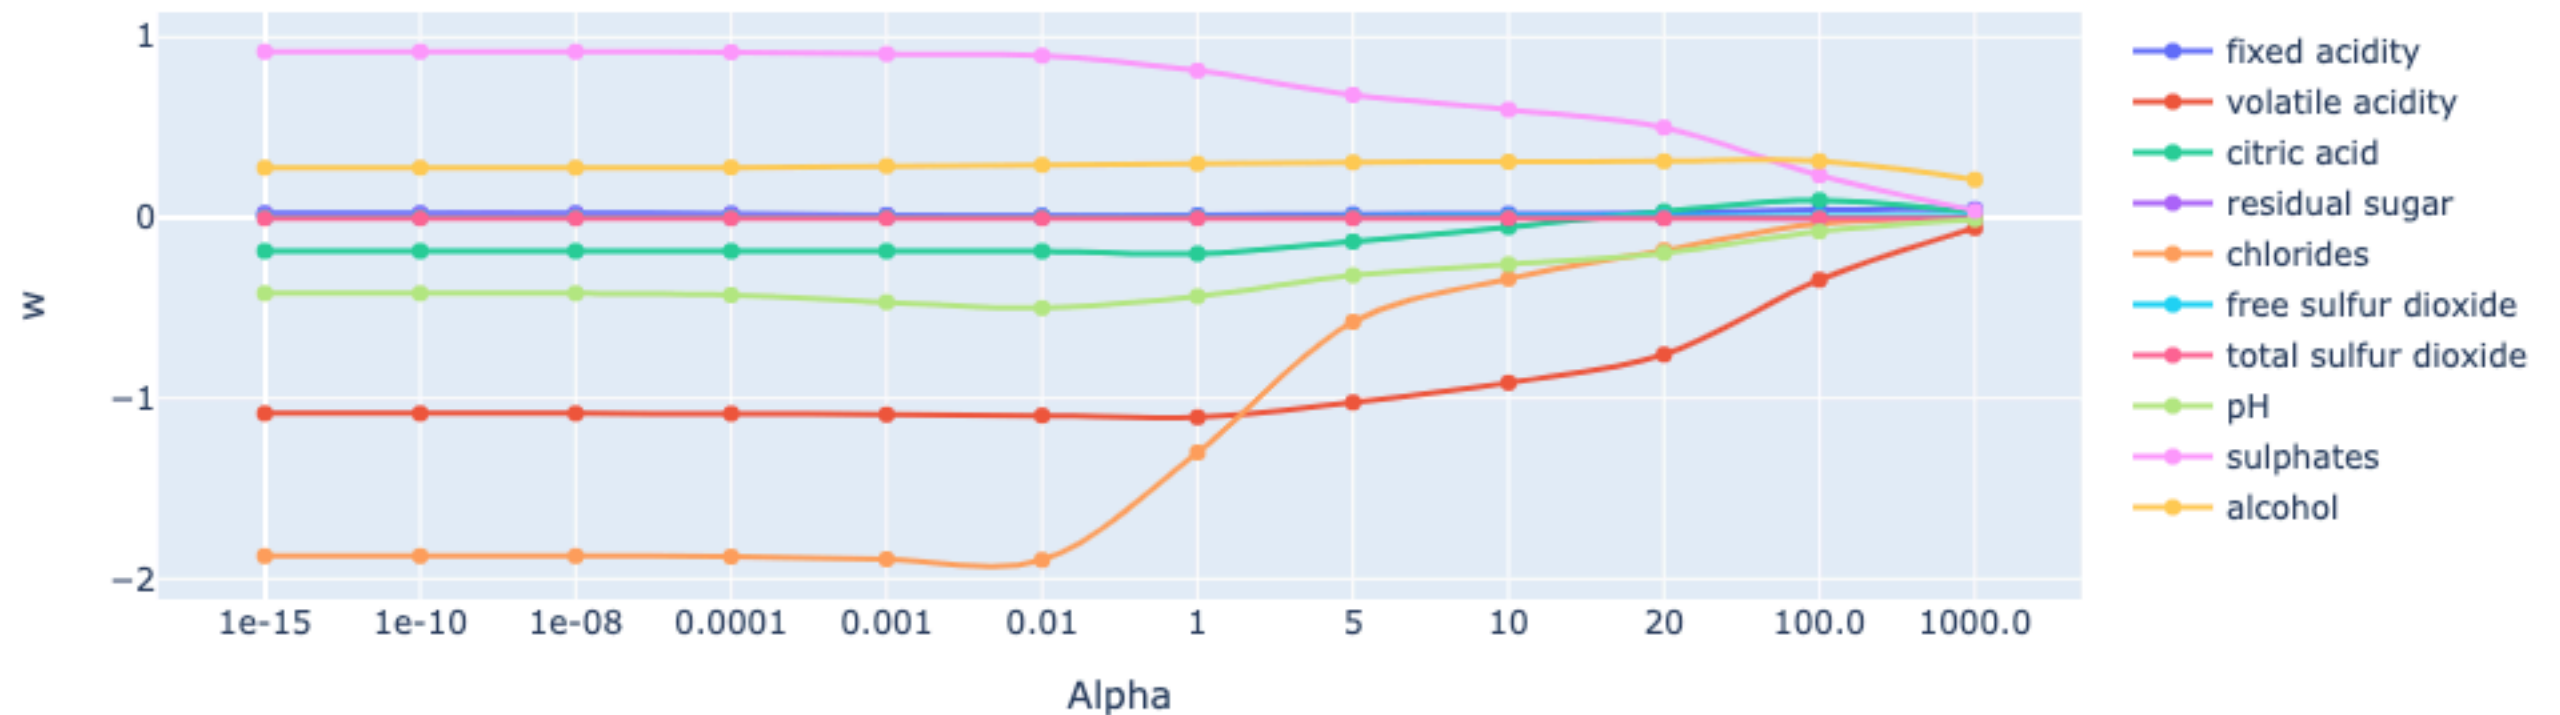
- Test Set

MSE	Adjusted R ²
0.434	0.2844

Regression

Multivariate linear regression with Ridge and Lasso Regularization (1)

- Ridge



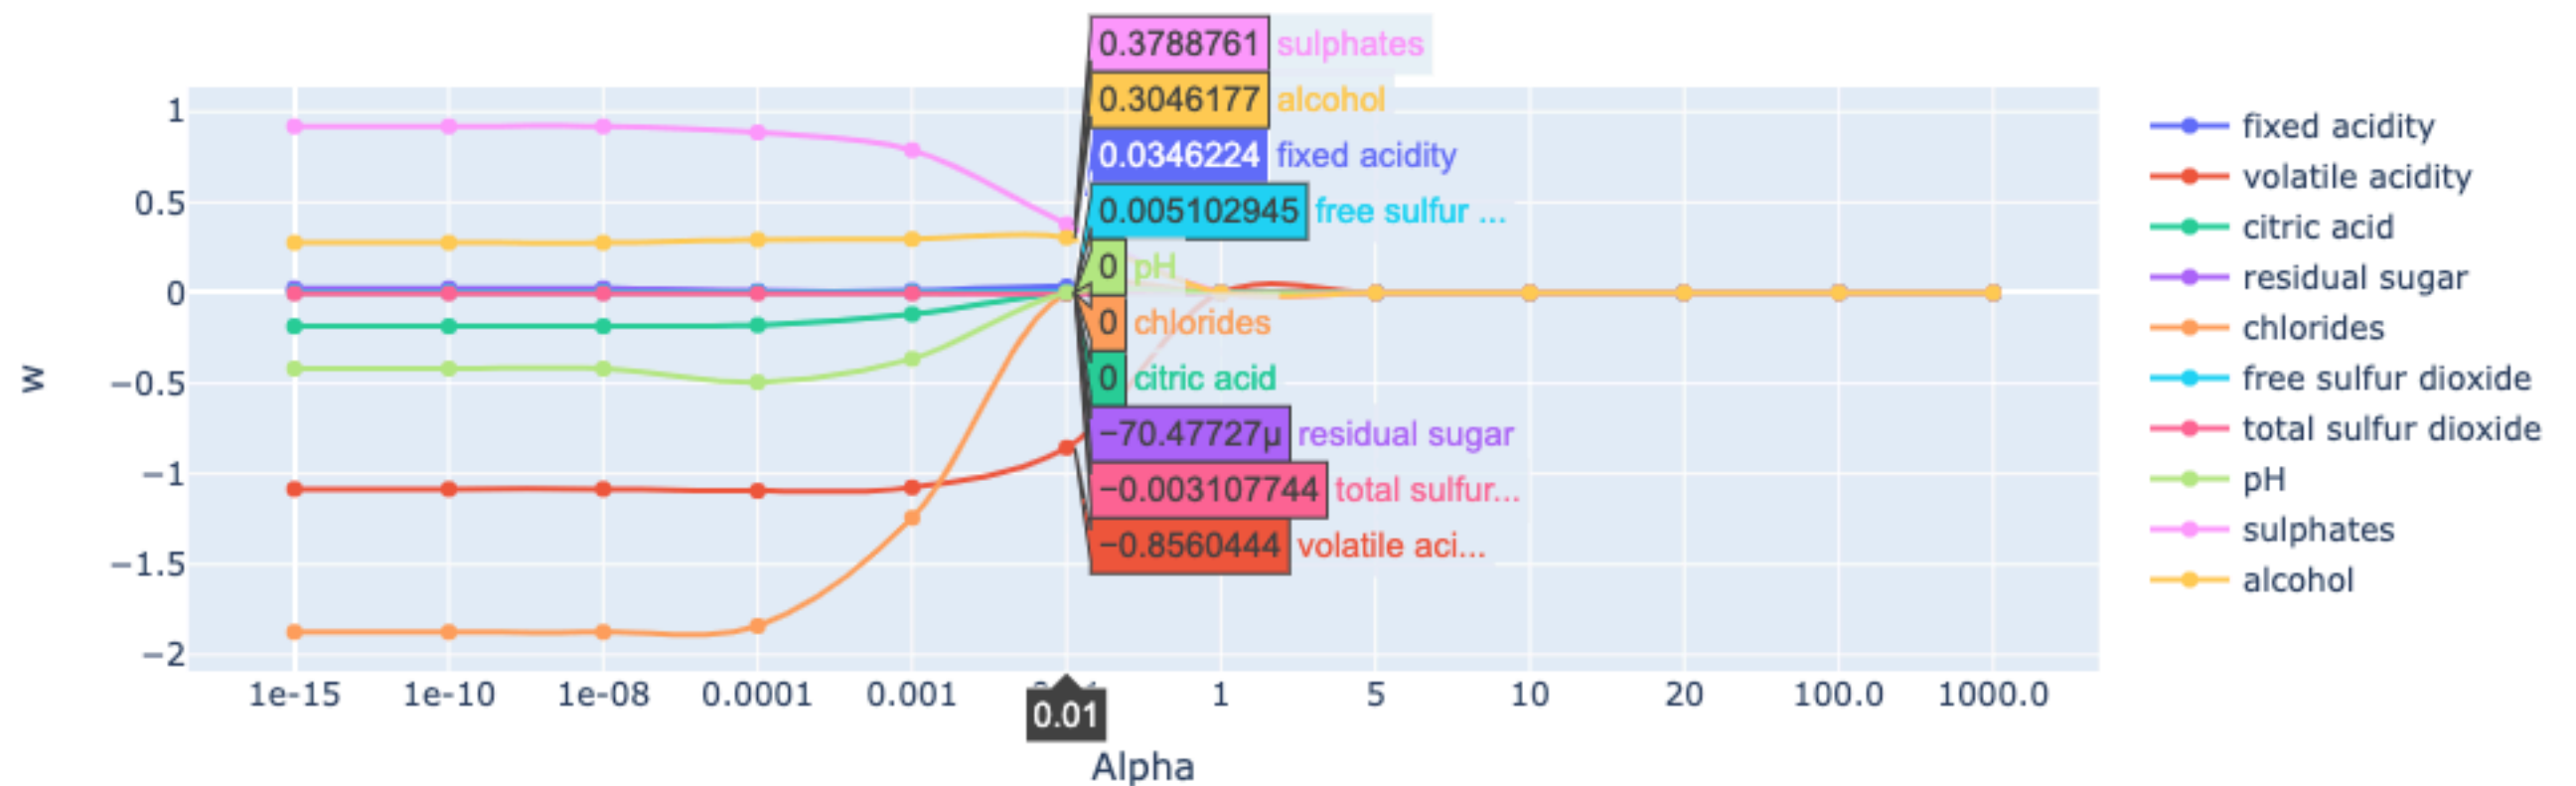
MSE
0.432

Adjusted R²
0.2867

Regression

Multivariate linear regression with Ridge and Lasso Regularization (2)

- Lasso


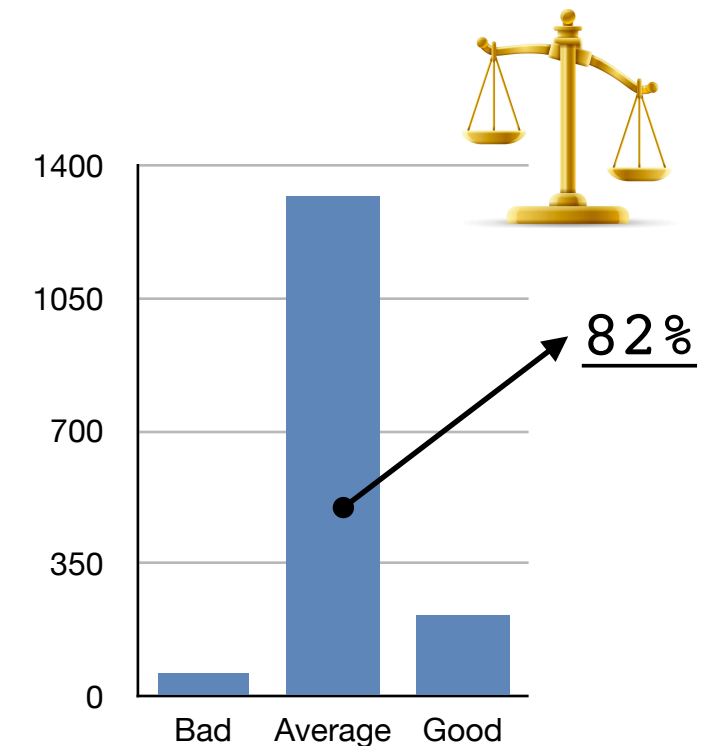


MSE
0.432

Adjusted R²
0.287

Classification

- Three classes: *Bad*, *Average* and *Good* wine
- Logistic Regression
K Neighbors Classifier
Support Vector Classifier (SVC)
Gaussian Naive Bayes
Gaussian Process Classifier
Decision Tree Classifier
Random Forest Classifier
- Playground:



Classification

Random Forest Classifier

```
> -----  
              PREDICTED  
            Bad  Average Good  
MEASURED Bad    0      13    0  
          Average 3     256    5  
          Good   0      23   20  
  
            precision    recall  f1-score   support  
  
         Bad           0.00      0.00      0.00         13  
    Average           0.88      0.97      0.92        264  
         Good           0.80      0.47      0.59         43  
  
 accuracy                   0.86        320  
 macro avg           0.56      0.48      0.50        320  
 weighted avg          0.83      0.86      0.84        320  
  
Misclassification cost: 44  
----- <
```

```
n_estimators = [100, 300, 500, 800]  
max_depth = [5, 8, 15, 25, 30]  
min_samples_split = [2, 5, 10, 15, 100]  
min_samples_leaf = [1, 2, 5, 10]  
random_state = [42]  
  
hyperF = dict(n_estimators = n_estimators,  
              max_depth = max_depth,  
              min_samples_split = min_samples_split,  
              min_samples_leaf = min_samples_leaf,  
              random_state = random_state)  
  
# Let's perform some tuning  
rf = RandomForestClassifier()  
grid_rf = GridSearchCV(estimator = rf, param_grid = hyperF, cv = 10,  
                       grid_rf.fit(xTrain, yTrain)  
bestParams = grid_rf.best_params_
```

Accuracy Score

0.863

Classification

Support Vector Classifier (SVC)

```
> -----
              PREDICTED
              Bad   Average Good
MEASURED Bad      0      13     0
          Average  0     261     3
          Good    0      25    18

              precision    recall  f1-score   support

         Bad            0.00      0.00      0.00         13
        Average         0.87      0.99      0.93        264
         Good            0.86      0.42      0.56         43

 accuracy              0.87         320
 macro avg              0.58      0.47      0.50        320
 weighted avg           0.84      0.87      0.84        320

Misclassification cost: 41
----- <
```

```
# Let's perform some tuning
model = SVC()
param = {
    'C': [0.1,0.8,0.9,1,1.1,1.2,1.3,1.4],
    'kernel':['linear', 'rbf'],
    'gamma': [0.1,0.8,0.9,1,1.1,1.2,1.3,1.4]
}
grid_svc = GridSearchCV(model, param_grid=param, scoring='accuracy')
grid_svc.fit(xTrain, yTrain) # ~ 1 min of tuning
bestParams = grid_svc.best_params_
```

Accuracy Score

0.872

Conclusion

- Quality entries not balanced
- Disappointing results but unrelated data
- Implementation of the topics covered in the course ✓
- Learning methods ✓

<https://github.com/nicorbtt/RedWineMachineLearning>